

Project 2B

CS111

02/15/19

What's new ?

- `lab2_list.c` -> `--lists` option
 - `--takes` number of lists as a parameter
 - Initialize several shared lists
 - Decreases contention !
- `profile.out` -> execution profiling report
 - How much time was spent in spin-locks and mutexes?
 - Add 8th column of data to your `.csv`
- New graphs
 - Throughput
- `gperftools` -> multi-threaded application friendly
 - Heap checker
 - Heap profiler
 - Cpu-profiler -> what we'll use

gperftools - Install

- Releases found :
<https://github.com/gperftools/gperftools/releases>
- Unpack and cd to directory
- `./configure --prefix=<path>`
- `make`
- `make install`
- `(make clean)`

Rmk: You do not need root permission to do this

Rmk: You also need to install `gv` to access graphical output, this is optional

gperftools - Use

The pprof command will allow you to analyse the profiling result.

To turn on CPU profiling, you have two options:

1. Define environment variable CPUPROFILE to the filename to dump the profile to (in your makefile)
Rmk: delete that file before each run
2. Bracket the code you want profiled using ProfilerStart() and ProfilerStop(). ProfilerStart() takes the profile filename as argument

Updating your makefile : profile option

1. Link the library:

`LD_PRELOAD = <path_to_lib>` (should be in `/usr/lib(64)`)

2. Link file to record dump

`CPUPROFILE = ./raw.gperf`

3. pprof with `--text`

`google-pprof --text <executable> <dump file> > $@`

4. pprof with `--list`

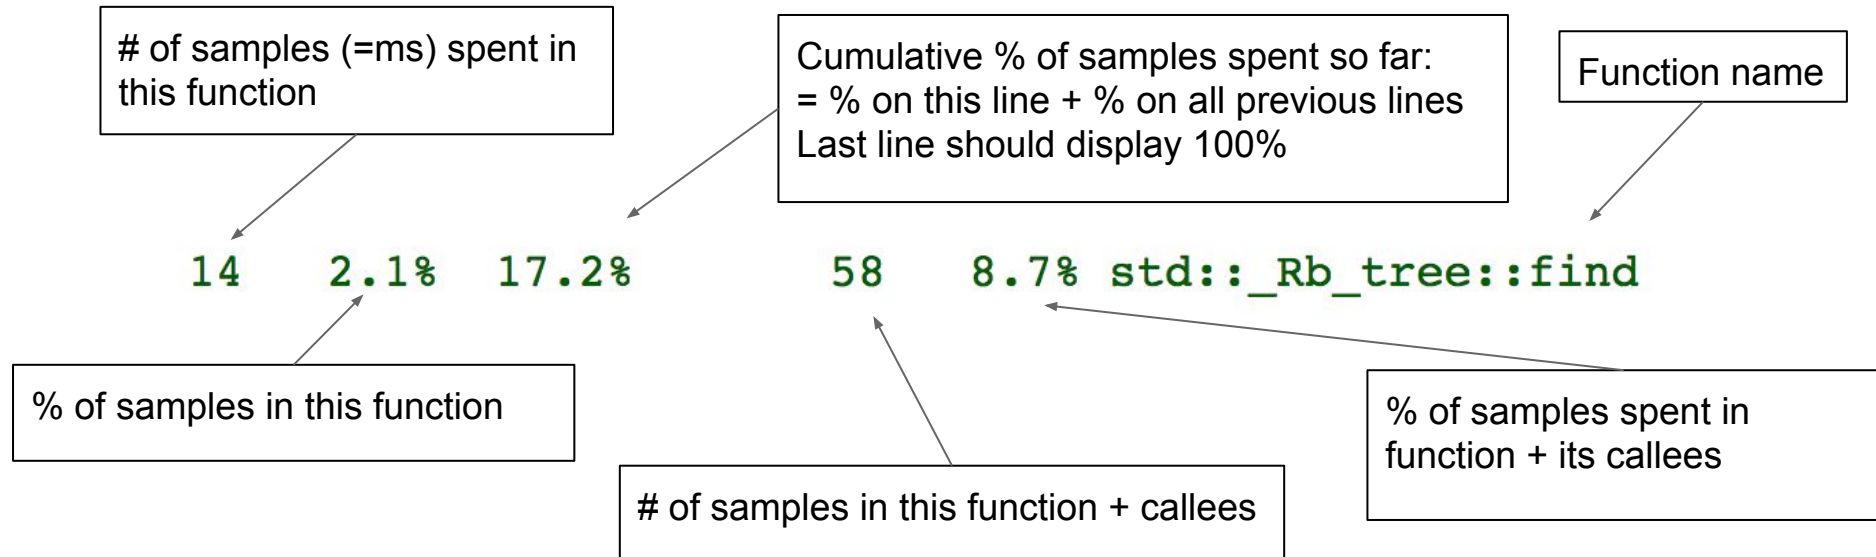
`google-pprof --list=<function> <exec> <dump> >> $@`

(\$@ is `profile.out` in this case)

Rmk: You also have to compile with the `-lprofile` flag

Analyzing the output

- [For more information](#)
- You **need** to use the text output, but if you first want to use the graphical output to make things clearer you can
- Text output will produce lines that look like:



Timing Mutex Waits

- Not spinning -> can't use profiling
- Instead : we will compute wait-for-lock time
 - Measure time before and after getting the lock
 - Add up all wait time for all threads
 - Divide by number of operations
 - Output Statistics for the run
- If you are not locking -> should get 0
- For example:
 - Allocate an array of timers (# threads)
 - `clock_gettime`, `pthread_mutex_lock([sublist])`, `clock_gettime`
 - add time to `timer[threadnum]`

Addressing the Underlying Problem

- Degradation is a result of increased contention
- To decrease contention : split the list
 - Add a lists option to specify the number of chunks
 - Select which list to insert node in
 - Get list length
 - Delete inserted nodes
 - exit
- To select which list : hash function
 - Maps threads*iterations -> #lists (key_value % lists)