

Project 3A

CS111 Discussion 1C

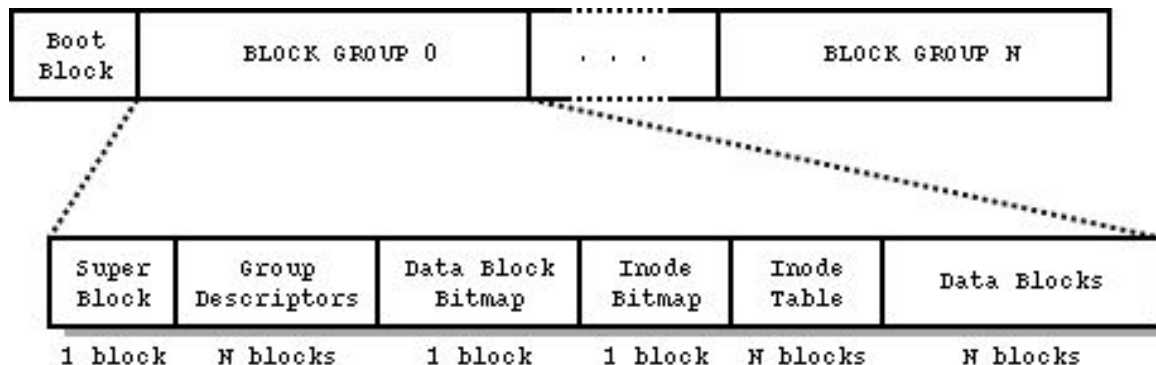
Note: You can form groups of 2 for this project

EXT2 filesystem

- You will be provided with a filesystem image (EXT2)
- You will explore it
 - You will mount it as a file, which means that:
 - You can read a specific number of byte at a specific offset
 - You will have to repeat these operations for all content
 - The problem becomes calculating the correct offset
- You will output a summary to stdout, describing:
 - The Superblock
 - Groups
 - Free Lists
 - Inodes
 - Indirect Blocks
 - Directories
- This involves a lot of simple code, **if** you have a clear image of what a filesystem is

EXT2 Filesystem

Read [this](#) carefully. Everything you need is in there.



Where to start? We want to find correct offsets.
The superblock! We know where it is

Note: File offsets are not addresses in memory! They can't be dereferenced with *

Note: A function that takes in a block number and outputs its offset would be useful

Superblock

- Contains all information about the filesystem:
 - # inodes
 - # blocks
 - How many blocks/inodes are free
 - # blocks/inodes per group
 - $\log(\text{block size})$
- Is located at an offset of 1024 bytes from the start of the device
 - This is the only data that we know how to localize, it is the starting point
 - We know the superblock is 1 block long
- You will need to refer to superblock data several times: make it accessible! (save it in a structure)

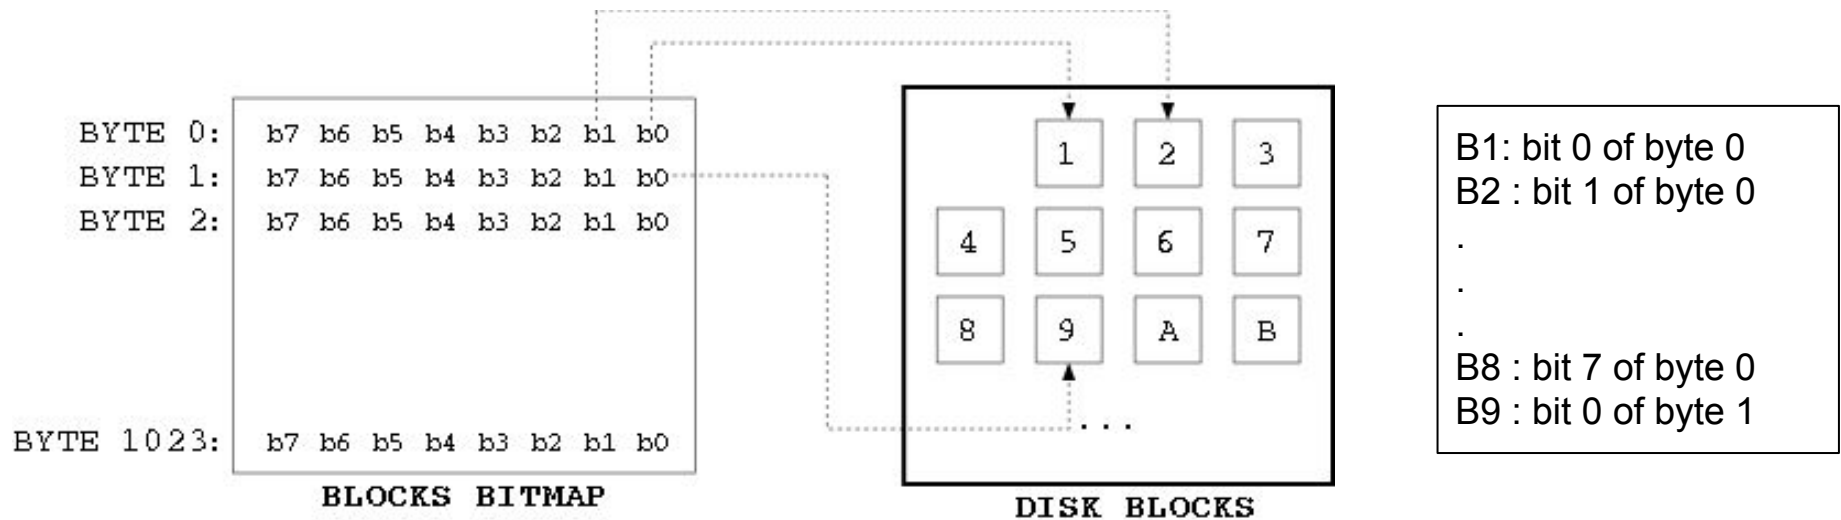
Block Group Descriptor Table

- Located at block 2, so at $\text{block_size} + 1024$
 - Block size is in the superblock
- It is an array of block group descriptors. Each of these block group descriptors provide (for 1 group):
 - The location of the inode bitmap/table
 - The location of the block bitmap
 - # of free blocks/inode
- This table can be stored on multiple blocks
 - How many? -> Superblock

Superblock offset

Block Bitmap

- Where is it located?
 - Bg_block_bitmap of its block group descriptor
- What does it represent?
 - Current state of a block within that block group (1 = used / 0=free)
- How is the information represented?



Inode Bitmap

- Where is it located?
 - Bg_inode_bitmap in group descriptor
- What does it represent?
 - Current state of an inode within the inode table (1=used)
- Works similarly to a Block Bitmap

Inode Table

- Where is it located?
 - `bg_inode_table` in the group descriptor
- It is an array of inodes. Each of those represent a single file (a directory, socket, buffer, symbolic link, regular file):
 - size and owner
 - access/modification/creation times (seconds since 1/1/70)
 - the number of blocks containing the data of the inode
 - which blocks contain the data of the file pointed to by the inode
 - Note that there is no filename in the inode
- The first few entries of the inode table are reserved

Locating an inode


- They are numerically ordered
 - Inode number \rightarrow index in the inode table
- The size of the inode table is fixed at format time
 - There is a cap on the number of entries
 - The size is (# of inodes * inode size) (cf superblock)
- Therefore the local inode index (relative to the current block group) is :
 - $(\text{inode_idx} - 1) \% \text{inodes_per_group}$

Inode pointer structure

- The inode points to the blocks that contain its data
- The block numbers are stored in a table
 - The first 12 are 'direct' blocks
 - The 13th block is the 'indirect block'
 - The 14th block is the 'doubly indirect block'
 - The 15th block is the 'triply indirect block'
 - If a 0 is encountered, all subsequent block pointers should be 0 (no other blocks defined)
- The indirect block structure is used to be able to reference files that would need more than 15 blocks with a single (fixed size inode)

Inode pointer structure

A picture is probably helpful:



Debugfs

- Download the image and mount it
 - This requires sudo: work on your machine
- You can explore it using debugfs(8)
 - You will need to research this tool to properly interpret values

```

debugfs: stat file1
Inode: 2790782   Type: regular   Mode:  0600   Flags: 0x0   Generation: 46520506
User:  2605    Group:  2601    Size: 14
File ACL: 0     Directory ACL: 0
Links: 1    Blockcount: 8
Fragment:  Address: 0    Number: 0    Size: 0
ctime: 0x3be712ea -- Mon Nov  5 15:30:02 2001
atime: 0x3be712ea -- Mon Nov  5 15:30:02 2001
mtime: 0x3be712ea -- Mon Nov  5 15:30:02 2001
BLOCKS:
5603924
TOTAL: 1
  
```

Dump content of inode

Access Control List

Inode change/access/modification time

Debugfs

- Other useful commands include:
 - `bd` : block dump
 - `testi <inode>` : test if inode is marked allocated in the inode bitmap
 - `testb <block>` : test if block is marked allocated in the block bitmap
- You can use this tool while writing your program to verify that you are exploring the filesystem correctly

General code structure

- Find/analyze/report superblock
- Copy info to common structure
 - Make the superblock data easily available to you
- Read group descriptions
 - How many blocks per group?
 - How many inodes per group?
- Find and report both bitmaps
- Analyze all allocated inodes
 - Is it a directory? A regular file?
 - What is the level of indirection?

Some useful notes

- Number of blocks per group?

Some useful notes

- Number of blocks per group?
 - You can find this in the superblock, except for the last group
- What is the block size?

Some useful notes

- Number of blocks per group?
 - You can find this in the superblock, except for the last group
- What is the block size?
 - $1024 \ll \text{superblock.s_log_block_size}$
- Is the file a regular file or a directory?

Some useful notes

- Number of blocks per group?
 - You can find this in the superblock, except for the last group
- What is the block size?
 - $1024 \ll \text{superblock.s_log_block_size}$
- Is the file a regular file or a directory?
 - Answer is in the first field of the inode : the 'mode'
- Write a function that takes in a block number and returns the absolute address (offset) of that block
- Write a function that reads a given number of blocks at a given offset

Iterating through inodes

- What is the level of indirection?
 - For a 1KiB block size:
 - Block 1 to 12 are direct
 - Block 13 to 268 are indirect
 - How many double indirect blocks?

Iterating through inodes

- What is the level of indirection?
 - For a 1KiB block size:
 - Block 1 to 12 are direct
 - Block 13 to 268 are indirect
 - How many double indirect blocks? $(256)*(256) = 65536$
 - How many triple indirect blocks?

Iterating through inodes

- What is the level of indirection?
 - For a 1KiB block size:
 - Block 1 to 12 are direct
 - Block 13 to 268 are indirect
 - How many double indirect blocks? $(256)*(256) = 65536$
 - How many triple indirect blocks? $(256)*(256)*(256) = 16777216$
- How does your offset change when you go through it?
 - Direct blocks -> blocksize
 - Indirect blocks?

Iterating through inodes

- What is the level of indirection?
 - For a 1KiB block size:
 - Block 1 to 12 are direct
 - Block 13 to 268 are indirect
 - How many double indirect blocks? $(256)*(256) = 65536$
 - How many triple indirect blocks? $(256)*(256)*(256) = 16777216$
- How does your offset change when you go through it?
 - Direct blocks -> blocksize
 - Indirect blocks? Multiply by blocksize per indirection level

A word on the trivial image

- You are provided with a small filesystem and the correct analysis output for that filesystem
- You are expected to be able to reproduce 'trivial.csv' from the given filesystem

BUT

- Doing so won't guarantee that your program is 100% correct, this tests basic functionality
- Also note that any access will modify the inodes!
 - This means that you won't be able to refer to the provided .csv
 - Use a fresh copy / mount in read only

FAQ

- How do I mount the image on a mac?
`hdiutil attach -readonly -nomount ./trivial.img`
- Can I assume that there will only be a single group in the filesystem used for grading?
Yes!