

# Project 1A

---

CS111  
Discussion 1A

# Breaking down project 1

- Instead of putting everything in a main, you may find it useful to write functions that:
  - Parses the next command and its arguments
    - Dynamically allocating memory for arguments
    - Checking that the syntax is correct
  - Executes a command given name and arguments
    - Creating the child process
    - Performing the necessary redirection
    - Executing the process
- Keep track of the number of open files
  - You can write a function that opens files/pipes
  - It increments a counter on the number of open files
  - It resets a dynamic flag variable after opening

# The verbose flag

- If this option is specified, print to stdout all options and arguments to that option before execution
  - To do this, write() to stdout, do not print()/putchar()...
  - Keep track of how many arguments have been passed to the current option
- getopt() will not recognise arguments, you will have to manually set optind as you iterate through them
  - getopt() will jump to the next option, which is a string starting with '--'

# fork(2)

Goal: Create a child process

**Success:**

- Parent receives child PID
- Child receives 0

**Error:** Parent receives -1

**pid\_t fork(void)**

```
#include<stdio.h>
#include<sys/types.h>
#include<unistd.h>
int main() {
```

What is the output ?

```
    fork();

    printf("Hello world!\n");
    return 0;
}
```

# fork(2)


## Goal: Create a child process

**Success:**

- Parent receives child PID
- Child receives 0

**Error:** Parent receives -1**pid\_t fork(void)**Important note:

When a parent calls fork(), the child that is spawned inherits all the open file descriptors from the parent! (Keep this in mind for later)

What is the output ?

```
fork();

printf("Hello world!\n");
return 0;

}
```

Hello world  
Hello world

2 processes running concurrently

# exec(3)

Goal: loading a new program in a calling process

Error : -1



```
int execvp(const char *file, char *const argv[]);
```



Path to new process image file

- Pointers to arguments available to new process image
- List is terminated by a NULL pointer

You know that:

- Every '--command' should be followed by at least 4 options (3 fd's + command name)
- Options for that '--command' are over when you're out of arguments or you see the next '--' in argc
- Creating a 'command' structure that holds all options for given commands could be useful
- Terminate it by NULL

# kill(2)

Goal: send a signal to a process (or a group)

Success: 0 // Error: -1

```
int kill(pid_t pid, int sig)
```

Receiving process id

Signal to be sent

(pid>0)? Send to process **pid**

(pid==0)? Send to all processes in group of sender

(pid==-1)? Send to all processes in group of calling process

(pid<-1)? Send to all processes whose gpid = abs(**pid**)

What signal will you send ? SIGINT for 0x03

# waitpid(2)

Goal: wait for child process to change state

pid of terminated child / **Error: -1**

Location to store exit information

**pid\_t** waitpid(**pid\_t** *pid*, **int** \**status*, **int** *options*)

- **<-1** : wait for any child process whose `gp_id == abs(pid)`
- **-1** : wait for any child process
- **0** : wait for any child process whose `gp_id == (calling process pid)`
- **>0** : Wait for the child whose `pid == pid`

- OR of these parameters (cf man):
- **WEXITSTATUS** : return exit status of the child
- **WIFSIGNALED** : returns True if child was terminated by a signal
- **WTERMSIG** : returns the number of the signal that caused the child to terminate



# waitpid(2)

When will you use it ?

- To harvest the exit status of a child process
- Or the signal number that caused the process to exit
- You will only wait on previous commands!
  - Store the pid of the commands you run through fork() calls
- If an exit status is available, print it out
- If the child was terminated by a signal, return the signal's value

# File flags

- There seems to be a lot of options in 1B...
  - But most of them are super simple!
- Create an open\_flag variable
  - Initialize it to 0
  - Reset it to 0 every time you open a file
- You can now implement all flags with an OR operation
  - `File_flag |= <option>`

# Pipe(2)



- Used for interprocess communication (IPC) on the same host
- When a call to `pipe()` is made, 2 file descriptors are returned
- One is the write end (on the left) and one is the read end (on the right)
- The idea is for process A to write(`fd_out`) and B to read (`fd_in`) : data goes from A to B!
- Note that pipes are unidirectional! If A writes to B, B **cannot** write to A. (need a 2nd pipe)

Should we call `fork()` before `pipe()` or after? Does it matter?

# Pipe(2)



- Used for interprocess communication (IPC) on the same host
- When a call to `pipe()` is made, 2 file descriptors are returned
- One is the write end (on the left) and one is the read end (on the right)
- The idea is for process A to write(`fd_out`) and B to read (`fd_in`) : data goes from A to B!
- Note that pipes are unidirectional! If A writes to B, B **cannot** write to A. (need a 2nd pipe)

Should we call `fork()` before `pipe()` or after? Does it matter?

- Important: Both processes A and B need to have access to the file descriptors associated to the pipe they will communicate with.
- Remember: when process A calls `fork()` to spawn process B, process B inherits all fds from A
- We need to call `pipe()` (creating 2 new fds in process A) and then `fork()` (so that B also gets access)
- If we `fork()` first, A and B don't have the same file descriptors anymore!

# Pipe(2) : child is the writer

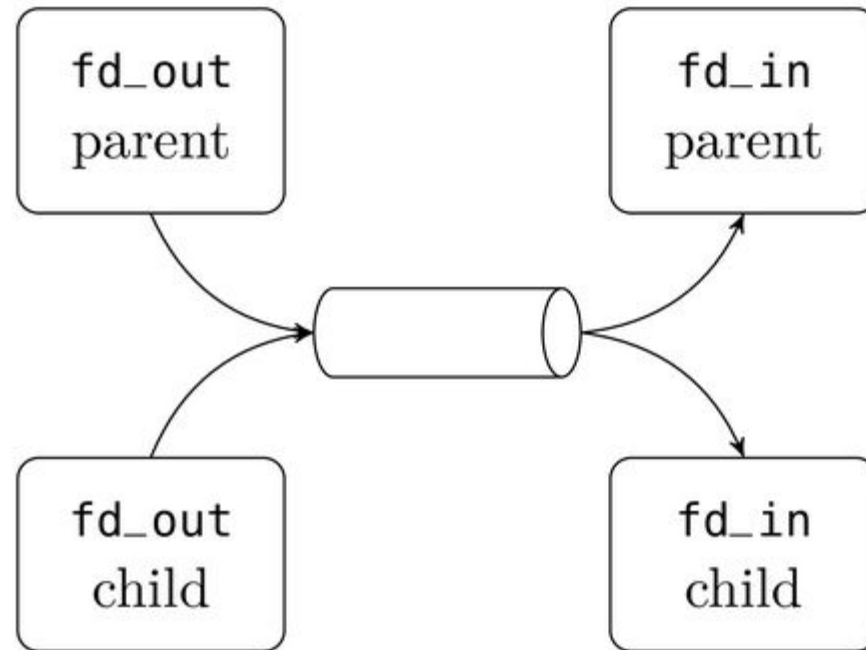
Senario: child writes to parent. Which file descriptors should we close?

Note:

This question arises because since we called `pipe()` and then `fork()`, both child and parent have access to both read and write ends of the pipes (4 fds total).


Pipes are unidirectional, so we will only use 2 of those 4. To pick the direction we have 2 fd's to close: one write end and one read end.

Closing unused file descriptors when piping is necessary (to harvest exit statuses)



# Pipe(2) : child is the writer

Close 1 fd on the parent side and  
one on the child side



# Pipe : close your file descriptors!

- In this project, you will create pipes between different children
- Just like in part A, you will have to do redirection after the call to `fork()`
  - To define where your child process takes input and where it outputs
  - In 1A, you did redirection with regular files
  - You can handle pipes in the same way!
- You will need to close the ends of the pipe you're not using
  - The simplest policy is to systematically close all file descriptors larger than 2 (in the child) **after redirection**

# --catch

- There are 3 signal options
  - You can use either `signal()` (simpler, deprecated) or `sigaction()` (more complete and complicated) to implement these
- For `--catch`, just register a signal handler
  - Print out the signal number and 'caught' to `stderr`
  - And `exit(signal number)`
- You can pass your signal handler to `signal()` or `sigaction`
  - Both calls just aim at defining what behavior your process should have when the specified signal arises
  - The second argument of `signal()` is the handler
  - For `sigaction` you have to define a 'sigaction' structure
    - Pass the handler through the `sa_handler` field



# --ignore

- Once again, you want to modify the behavior of your process if it receives a given signal
- You are doing the same thing as in --catch, the only thing that changes is the handler you specify
- In catch, it was one of your functions
  - Now you have to ignore the signal
- SIG\_IGN is the disposition that should be passed as a handler
  - Second argument of signal()
  - The second argument to sigaction() is a pointer to the sigaction structure, set its sa\_handler field to SIG\_IGN

# --default

- We now want the 'default' handler
- SIG\_DFL

# --abort

- Force a segfault
- The system should dump core via a segmentation violation
- Just like lab0! Dereference a null pointer

# --pause()

- pause , waiting for a signal to arrive
- The C standard library provides a function pause()
  - It causes the the calling process to sleep until a signal is delivered
- Just call it!

# getrusage()

Goal: Measure resource usage.

Fills out this struct:

```
struct rusage {
    struct timeval ru_utime; /* user CPU time used */
    struct timeval ru_stime; /* system CPU time used */
    long ru_maxrss; /* maximum resident set size */
long ru_ixrss; /* integral shared memory size */
long ru_idrss; /* integral unshared data size */
long ru_isrss; /* integral unshared stack size */
    long ru_minflt; /* page reclaims (soft page faults) */
    long ru_majflt; /* page faults (hard page faults) */
long ru_nswap; /* swaps */
    long ru_inblock; /* block input operations */
    long ru_oublock; /* block output operations */
long ru_msgsnd; /* IPC messages sent */
long ru_msgrcv; /* IPC messages received */
long ru_nsignals; /* signals received */
    long ru_nvcsw; /* voluntary context switches */
    long ru_nivcsw; /* involuntary context switches */
};
```

# timeval

- The measured time is stored in a structure, which has 2 fields:
  - tv\_sec : number of **whole** seconds of elapsed time
  - tv\_usec : rest of the time in microseconds
- $\text{Time\_elapsed} = \text{tv\_sec} + \text{tv\_usec}$ 
  - The 2 fields are not the same value in different units!

# getrusage()

- You have to specify “who’s” resource usage you want to monitor
- You want to monitor every option
- The options are:
  - RUSAGE\_SELF : monitor calling process (sum of all threads)
  - RUSAGE\_CHILDREN : monitor all descendants that have terminated or have been waited for
  - RUSAGE\_THREAD : monitor the calling thread
- Where will you use it?
  - When using catch, close, pipe, opening a file -> just the calling process
  - When executing a command -> also monitor the child process!